

Functorial Type Theory: An Executable Architecture for Directed Dependency

The emdash contributors

overview research article

Research-draft status. This article describes the checked emdash v3.2 development and a bounded TypeScript reviewer. It is not a released foundation, a venue submission, or a claim that the full research programme is complete. The active Lambdapi sources remain the mathematical authority.

Companion book [8]: [10.5281/zenodo.21544186](https://zenodo.org/record/21544186). Code [4]: github.com/hotdocx/emdash. The published book record and ongoing source may carry different versions.

Abstract

Dependent type theory makes substitution computational. Category theory makes variation along arrows explicit. **Functorial Type Theory** asks what happens when both principles inhabit one language: variables may range over objects and arrows, dependent families carry transport, and functoriality and naturality are internal operations that can normalize rather than external proof obligations.

Emdash is an executable investigation of this idea. Its active Lambdapi development contains an autonomous categorical or directed dependent type theory of categories, Cat-valued families, reindexing, total categories, section categories, dependent homs, functors, and transfors. Lambdapi's dependent $\lambda\Pi$ framework supplies a complementary outer logical layer. A small TypeScript implementation mirrors this separation: an explicit locally nameless dependent Core checks and evaluates terms, while a typed categorical frontend compiles usable binder notation to existing internal categorical owners, and bounded typed declarations expand into ordinary logical-framework declarations and rules.

Five computations summarize the architecture. First, for $p : x \rightarrow y$ and $q : y \rightarrow z$, emdash forms the

This overview concentrates on the foundational architecture. The companion book also develops the newer monad and Cartesian/indexed interfaces, whole kernel/cokernel and homology operations, and retained proof-CAS computation. The current whole-opposite and Sigma-Hom encodings have known soundness defects; higher-variance-dependent formal claims remain unqualified pending repair. Local checks in the research artifact are not a consistency certificate.

outgoing-arrow category

$$\text{PathOut}_Z(x) = \Sigma_{y:Z} \text{Hom}_Z(x, y)$$

and a canonical arrow $\rho_{x,y,p} : (x, \text{id}_x) \rightarrow (y, p)$. Transporting a motive along ρ gives a synthetic arrow-induction principle. Applied to the representable composition motive, its checked normal form is $q \circ p$. Second, for a ring map out of R , the condition that $f \in R$ becomes invertible defines an ordinary sieve $D_R(f)$ before any representing open is chosen. A supplied localization represents that sieve pointwise, finite unit-ideal families generate the big Zariski topology, and a separate direct cover completion constructs a Cat-valued sheafification reflector from return, cover-indexed glue, and silent coherence. Third, localizing natural numbers along successor gives an integer line, and an opaque groupoidal Circle HIT carries a universal Integer cover. Encode/decode proves $\text{Hom}_{S^1}(\text{base}, \text{base}) \simeq \mathbb{Z}$; the directed walking endomorphism maps to the Circle by the nonnegative inclusion, while product paths admit both coherent sequential transport factorizations. More generally, $\text{Groupoidify}(C)$ is characterized by a whole equivalence between groupoidal maps out and path-

valued functors on $\mathcal{C}$. A selected strict-object/lax-arrow Gray closure then recovers a nonidentity walking-square interchanger from the same internal laxity action. Fourth, face codes and directed join supply ordinal simplexes, while iterated outgoing paths give their dependent cells. One Nat recursion constructs their canonical source in variable dimension, with faces and higher action checked through dimension four. Fifth, the TypeScript frontend accepts ordinary, natural, displayed-functorial, and displayed-natural abstractions. It recursively factors variable occurrences through weakening, pairing, evaluation, reindexing, totalization, and internal action owners, then emits backend-neutral explicit Core. The same checker accepts the result at

object and arrow level; unsupported factorizations fail closed.

The result is a working research artifact rather than a completed proof assistant. It demonstrates a coherent design across mathematical kernel, dependent logical framework, elaboration, checked computation, and a client-side reviewer. Arbitrary dependency and variance graphs, whole-library transfer, systematic groupoidal closure for every former, source-functorial adjunction packaging for groupoidification, full Gray monoidality, a commutative-ring lift and left-exactness theorem for constructed sheafification, representation-independent schemes, and global metatheory remain explicit research boundaries.

1. Why Categorical Variables Should Compute

In ordinary dependent type theory, a bound variable may occur anywhere in a term, and the elaborator reconstructs applications and substitutions from local typing information. In ordinary category-theoretic prose, the same convenience is assumed for a richer notion of variable. A symbol may denote an object, an arrow, a functor, or a natural family; an expression such as $F(x)$ silently invokes the appropriate object or arrow action. Naturality is usually discharged by saying that a construction is "functorial in x ."

That phrase conceals an implementation choice. One can represent a construction as a pointwise function plus externally supplied equations, or one can represent it by an internal functorial object whose action and higher action are already part of the term. Emdash chooses the latter whenever a stable categorical construction is available. The aim is not to attach a naturality proof to every pointwise definition. It is to elaborate readable syntax into compositions of internal owners for which naturality follows from the type and generic action calculus.

This makes normalization do mathematical work. A composite of represented-hom actions is a cut; a rewrite that contracts it is cut elimination. A transformation component is not merely a projection from an opaque proof of naturality; it is an observable rung of an iterable hom tower. A dependent total arrow is not only a pair of endpoints; it contains a base arrow and a fibre arrow above transport.

The project therefore uses *dependent type theory* in two complementary senses.

1. An **outer dependent logical framework** supplies Π -types, λ -abstraction, application, definitions, and conversion. Lambdapi provides this layer for the authoritative development. The TypeScript Core implements a bounded counterpart with dependent checking and β/δ computation.
2. An **inner categorical or directed dependent type theory** treats a category as a context or shape and a Cat-valued functor as a dependent category over it. Reindexing is substitution, a total category is dependent sum, a section category is dependent product, and dependent hom records transport over a directed base arrow.

Neither layer replaces the other. The outer framework hosts the inner calculus; the inner calculus gives computational meaning to variables whose variation is genuinely categorical. The current groupoidal/type-theoretic universe supplies equality, J , dependent pairs, and dependent products as well. A checked Circle/Integer encode--decode theorem and representative product-transport closure connect these layers. Category-indexed groupoidification now adds the target-side universal mapping property, while its source action and packaged adjunction—and systematic closure for every former—remain later work.

The same order of construction extends beyond binder syntax. A local condition should first be stable under every change of stage; only afterward should one ask

whether one object represents all successful probes. Thus the condition that a section becomes invertible is first an ordinary sieve. In affine geometry a localization may represent it, and finite families of such charts may generate a topology. Likewise, sheafhood is first a property of restriction from whole sections to whole matching families; a sheafification reflector is a further construction. These distinctions let categorical semantics carry computation without importing an abstract modal layer or silently assuming the classical existence theorems suggested by familiar notation.

The central claim of this overview is deliberately narrower than a foundational completeness theorem:

A substantial directed dependent calculus, a sieve-to-sheafification application, a computational

groupoidal-realization slice, a minimal outer dependent framework, and a recursive categorical-binder frontend already compose into one executable architecture.

The rest of the paper makes that claim concrete. Sections 2 and 3 describe the two layers and the directed dependent constructors. Sections 4 and 5 derive synthetic arrow induction and its composition normal form. Section 6 explains how ordinary and displayed categorical binders compile. Section 7 follows the result through explicit Core to the browser artifact. Sections 8 and 9 state the local-to-global geometric application, free groupoidal realization, and the broader programme's present limits.

2. A Two-Layer Executable Architecture

The authoritative implementation is a `Lambdapi` signature. It declares the categorical owners and orients selected structural equations as rewrite rules. Proof-time unification rules relate useful presentations when neither should become the runtime normal form. Executable assertions in a separate checks module exercise formation, reduction, comparison, and non-collapse cases.

The `TypeScript` implementation is not a second mathematical authority. It is a small product kernel and

elaboration target whose vocabulary is aligned with reviewed `Lambdapi` owners. Its trusted representation is explicit Core: applications have visible spines, binders are locally nameless, plicity and variation metadata are retained, and terms can be traversed, frozen, serialized, and compared structurally. A one-shot scoped builder restores HOAS-like authoring ergonomics without storing JavaScript closures in Core.

The two input paths meet at one generic checker:

Input	Deterministic expansion before checking
text or typed <code>TypeScript</code> terms	contextual occurrence analysis, typed categorical wiring, and backend-neutral explicit emdash Core
typed host declarations	validated expansion to ordinary LF declarations and rules

Both paths are consumed by the small `TypeScript` dependent-LF checker and bounded evaluator. Selected judgments may then be emitted to `Lambdapi` as a deterministic conformance oracle; the active `Lambdapi` v3.2 source remains the mathematical authority.

The outer `TypeScript` LF includes dependent Π , annotated λ -terms, application, contextual metavariables, transparent definitions, and bounded β/δ conversion. It deliberately rejects `Type` :

`Type` ; object theories use decoded codes in the same style as the `Lambdapi` development. The categorical layer does not add a new checker for each owner. Reviewed declarations and runtime or proof-time rules are compiled into the generic environment, after which ordinary LF inference, checking, conversion, and rewriting process the term.

The declaration branch removes repetition in the outer framework rather than adding categorical semantics. One convenience packages already typed adjunction

data, or a counit and whole hom transpose, as an indexed assumption with proof-time agreements. A second expands one unparameterized, nonrecursive, single-constructor dependent structure into an opaque carrier, constructor, primitive projections, and ordered

2.1 The outer dependent LF

The outer Core has only the generic syntax needed to host object theories: sorts, free and bound variables, dependent products, lambdas, applications, metavariables, and owner calls. A dependent product stores the domain and a body under one binder. A lambda stores the same binder information and an explicit body. Bound occurrences use indices, so alpha-equivalence is structural and substitution is a total traversal rather than an invocation of an opaque host-language function.

Inference and checking are bidirectional. A free owner is looked up in the candidate environment; a call consumes its dependent product spine, instantiating each codomain with the checked argument. A lambda is normally checked against an expected product. An annotated lambda can also infer a product, which is useful for direct construction and diagnostics. Implicit parameters remain visible in Core even when a surface builder inserts them.

Conversion is deliberately bounded and evidence-producing. Weak-head β instantiates a lambda body and preserves any remaining application spine. Authorized transparent declarations unfold by δ only after their bodies have been checked against their declared types and their dependency order has been validated. Metavariables are zonked before rigid comparison, and Miller-pattern constraints are revisited as information arrives. One shared budget prevents a

projection beta rules. Both forms disappear before ordinary Core checking. Neither adds a trusted term node, a runtime alias, general record eta or elimination, or a new Lambdaapi owner.

conversion query from hiding an unbounded reduction path; traces record whether progress came from β , δ , metavariable resolution, or a reviewed categorical runtime rule.

This is enough to call the layer a minimal dependent logical framework, but not enough to claim a new general-purpose foundation. It does not implement an unrestricted user rewrite language, assume **Type** : **Type**, or silently treat every declaration as transparent. Those exclusions are part of the trusted boundary. The original TypeScript prototype had more globally mutable conveniences, including holes, rewriting, and HOAS storage; the renewed Core keeps the reusable algorithms while replacing global authority and opaque binder bodies with explicit, profile-scoped data.

The scoped builder reconciles the explicit representation with an ergonomic API:

```
pi("x", A, x => B(x))
lam("x", A, x => body(x))
let_("x", value, x => body(x))
```

Each callback runs once. Its token is valid only during lowering and becomes a bound index immediately. A foreign or escaped token is rejected. Thus the authoring surface can look higher-order while the term that enters checking is first-order, deterministic, and serializable.

2.2 Semantic owners and transfer

The inner theory enters TypeScript through a typed immutable transfer representation. A declaration records its qualified owner, binder telescope, result, visibility, optional body, and source evidence. Runtime rules and proof-time comparisons are separate data classes. A policy overlay decides which reviewed items are executable in a candidate profile; it is not inferred merely because a Lambdaapi symbol can be named.

That distinction mirrors the mathematical source. An opaque declaration extends typing but cannot unfold. A transparent definition participates in δ conversion. A runtime rule selects an operational normal form. A proof-time rule assists a rigid comparison without orienting runtime evaluation. Compiled module interfaces preserve public, protected, and private visibility rather than flattening every source declaration into one global namespace.

The transferred profile is a dependency-closed selection rather than the entire library. Its generic engines have handled opaque and transparent declarations, grouped runtime rules, proof comparisons, source-ordered modules, generated inductive owners, internal Pi, Sigma-transfor operations, and profunctor fragments. A categorical operation is represented by a semantic owner and argument schema, not a bespoke TypeScript AST tag with a private evaluator branch. Thus, after mathematical review, extending the profile is principally a typed-data and policy operation. This does not imply that every remaining declaration can be imported as one batch or that a parseable rewrite rule automatically acquires authority.

The public text adapter is intentionally later and thinner. It recognizes a bounded mathematical syntax and

resolves it into the same typed categorical surface used by direct TypeScript construction. Parsing a string does not create semantics. Expected types, classifier information, and the available internal owners decide whether an application denotes an object action, an arrow action, a displayed section component, a whole hom action, or a supported higher component. The adapter fails with a source location when no internal factorization exists.

This is also why Lambdapi remains useful after a browser-capable TypeScript checker exists. The TypeScript runtime makes the demonstrated profile small, inspectable, and client-side. Deterministic Lambdapi emission compares selected judgments against the active source during development. The oracle is a conformance route, not a production backend and not a substitute for the TypeScript checker.

3. Directed Families, Totals, And Sections

Let K be a category. A directed family over K is a functor $E : K \vdash \mathbf{Cat}$. It supplies a fibre category $E[k]$ for every object k , but also transport along every base arrow:

$$p : k \rightarrow k' \quad \Longrightarrow \quad E[p] : E[k] \vdash E[k'].$$

The arrow action is essential. A pointwise assignment of categories is not yet a directed dependent type. In the kernel, the stable category of such families is `Catd_cat K`; morphisms and higher morphisms are exposed by `Functord_cat` and `Transfd_cat`. Their ordinary Cat-valued presentations are available through controlled projections and proof-time comparisons, but the stable displayed heads are retained so that higher action remains iterable.

Reindexing along $F : A \vdash K$ is substitution:

$$F^* E[a] = E[F[a]],$$

implemented by `Pullback_catd`. The total category is the directed dependent sum:

$$\Sigma_K E.$$

Its objects are pairs (k, u) with $u \in E[k]$. An arrow $(k, u) \rightarrow (k', u')$ consists of

$$p : k \rightarrow k', \quad \alpha : E[p](u) \rightarrow u'.$$

Thus the hom of a total category is organized by a dependent hom, not by an ordinary product detached from transport. The owner `sigma_arrow` forms (p, α) , and `sigma_transport_arrow(E, p, u)` uses the identity fibre component to produce

$$(k, u) \longrightarrow (k', E[p](u)).$$

The dependent product is the category of sections:

$$\Pi_K E.$$

A section s assigns $s[k] \in E[k]$ and carries an action

$$s[p] : E[p](s[k]) \rightarrow s[k'].$$

The stable `Pi_cat` facade exposes section objects and the next displayed hom. Evaluation and section action project through the generic displayed application tower. They are not a second primitive function calculus.

3.1 Dependent hom as shared infrastructure

The common shape behind total arrows and section action is the dependent hom.

Given

$$u \in E[x], \quad v \in E[y], \quad p : x \rightarrow y,$$

the fibre component of a directed arrow is

$$\text{Hom}_{E[y]}(E[p](u), v).$$

As p varies this expression must itself carry action. Emdash packages it as a **Cat**-valued construction over the base hom, with variance chosen so that higher pre- and postcomposition normalize in the intended direction. The public `homd` and internalized `homd_int` owners are therefore more than notation for a pointwise hom formula. They retain the object action, base arrow action, and the next hom needed by later consumers.

This internalization avoids a common design trap. Suppose one defined a displayed functor only by maps on fibre objects and fibre arrows and then attached an external naturality square. Such a package would be awkward to iterate: the next categorical dimension would have to unpack the square and rebuild its action. In emdash, the relevant object is already a `Functor_d_cat` or `Transfd_cat` term. Its action is observed through generic `fapp*`, `tapp*`, and displayed internal-hom projections. Higher consumers receive a first-class categorical term rather than a pointwise function plus evidence.

The same principle explains several otherwise technical owners. `sigma_map_func` turns a displayed functor into a functor between total categories, with its fibre action ending in the displayed internal-hom projection ladder. `sigma_pullback_total_func` maps the total of a reindexed family back to the original total:

$$\Sigma_A(F^*D) \longrightarrow \Sigma_K D, \quad (a, u) \longmapsto (F[a], u).$$

Its arrow action sends (p, α) to $(F[p], \alpha)$. The construction is asymmetric because one input is a family reindexed along a specified functor; it is not a generic pullback of arbitrary total functors.

Likewise, section evaluation at an object and section action along an arrow are related observations of one coherent section. The fixed-index `piapp0` view is useful at the surface, while the generic displayed application tower retains its arrow and higher action. This is why the directed calculus deserves to be read as a DTT in its own terms: substitution, context extension, dependent programs, and their directed action are part of one computational structure.

These operations iterate to genuine dependent telescopes. If $R : K \vdash \mathbf{Cat}$, then $\Sigma_K R$ is the extended directed context. A second family depending on both k and $r \in R[k]$ becomes a family over this total category. The owners `Sigma_catd_functor_d_catd` and `Sigma_transfd_funcd` internalize the corresponding family and transformation uncurrying. Further `Sigma_cat` and `Pi_cat` applications extend the telescope or form its dependent programs.

Independent displayed variables over one base use fibrewise products without postulating a new primitive product-family head. For $B, C : K \vdash \mathbf{Cat}$, write

$$P(B, C)[k] = B[k] \times C[k].$$

The active presentation is built from ordinary product formation, uncurrying, and the displayed family structure. Its projections, pairing, swap, and diagonal support weakening, exchange, and contraction among independent fibre siblings. This does **not** license swapping variables across a genuine dependency edge. In a telescope

```
k : K;
a : A[k];
b : B[(k,a)], c : C[(k,a)];
d : D[((k,a),(b,c))]
```

b and c share a base and may be paired or exchanged; a cannot be moved past data whose type depends on it. The distinction is the familiar structural discipline of dependent type theory, now interpreted over objects and arrows of a category.

4. Synthetic Arrow Induction

The mathematical anchor of v3.2 is a directed analogue of path induction.

Fix x in a category Z . The covariant representable family is

$$\text{Rep}_Z(x)[y] = \text{Hom}_Z(x, y).$$

Its total category is the category of outgoing arrows:

$$\text{PathOut}_Z(x) = \Sigma_{y:Z} \text{Hom}_Z(x, y).$$

An object is a pair (y, p) with $p : x \rightarrow y$. The distinguished object is the reflexive arrow

$$\iota_x = (x, \text{id}_x).$$

Every outgoing arrow has a canonical total arrow from ι_x :

$$\rho_{x,y,p} : \iota_x \longrightarrow (y, p).$$

It is not postulated as a new path constructor. It is ordinary Sigma transport in the representable family:

$$\rho_{x,y,p} = \text{sigmaTransport}(\text{Rep}_Z(x), p, \text{id}_x).$$

The endpoint is correct because representable action computes:

$$\text{Rep}_Z(x)[p](\text{id}_x) = p.$$

$$\begin{array}{ccc} (x, \text{id}_x) & \xrightarrow{\rho_{x,y,p}} & (y, p) \\ \pi_1 \downarrow & & \downarrow \pi_1 \\ x & \xrightarrow{p} & y \\ \text{fibre} \downarrow & & \downarrow \text{fibre} \\ \text{id}_x & \xrightarrow{\text{Rep}_Z(x)[p]} & p \end{array}$$

Now let $E : \text{PathOut}_Z(x) \vdash \mathbf{Cat}$ be a motive and $u \in E[\iota_x]$. Transport along the canonical arrow defines a section:

$$\text{Ind}_x(E, u)[(y, p)] = E[\rho_{x,y,p}](u).$$

In kernel notation, `path_ind_sec(Z, x, E, u)` inhabits

$$\prod_{q:\text{PathOut}_Z(x)} E[q].$$

This is fixed-source arrow induction. It resembles identity elimination because the reflexive outgoing arrow determines all components, but no invertibility is assumed. The comparison from reflexivity to (y, p) is a directed arrow in a Sigma category.

The principal theorem internalizes the source x as well. Precomposition by $r : x \rightarrow y$ gives

$$\text{PathOut}_Z(r) : \text{PathOut}_Z(y) \vdash \text{PathOut}_Z(x).$$

$$\text{PathOut}_Z(r)[z, q] = (z, q \circ r).$$

Consequently motives vary by pullback and sections vary by section pullback. The source-indexed induction theorem is a displayed transformation

```
PathInd(Z) :
  PathOutReflEval_Z => PathOutPi_Z
```

Its kernel owner is `PathInd_transfd(Z)`. At each x it is the fixed-source induction functor. Along an arrow r , its type already contains the comparison between transporting initial data and pulling back the resulting section. Naturality is therefore internal to one displayed transformation; it is not an externally supplied square.

The fixed-source component can be pictured as motive transport:

$$\begin{array}{ccccc}
 E[(x, \text{id}_x)] & \xrightarrow{\quad} & E[\rho_{x,y,p}] & \xrightarrow{\quad} & E[(y, p)] \\
 \vdots & & & & \vdots \\
 \text{element} & & & & \text{element} \\
 \downarrow & & & & \downarrow \\
 u & \xrightarrow{\quad} & \text{Ind}_x(E, u)[(y, p)] & \xrightarrow{\quad} & E[\rho_{x,y,p}](u)
 \end{array}$$

There are two different kinds of non-strictness here. First, changing the source object acts contravariantly on `PathOut`, so the target section is pulled back rather than transported by an identity. Second, an arbitrary displayed functor need not preserve a transported Sigma arrow as a strict cartesian map; its component-level comparison remains visible through the displayed internal-hom action. The theorem packages the comparisons that are actually justified. It does not obtain a stronger strictness principle by calling the construction "induction."

There is also a totalized presentation over the Sigma category of sources and motives. It is derived by applying `Sigma_transfd_funcd` to the telescope theorem:

```
PathInd_funcd(Z) =
  Sigma_transfd_funcd(PathInd_transfd(Z)).
```

The direction of ownership is significant. The displayed transformation is primary because it retains the varying-source action. The total functor is a useful derived view, not a second induction axiom.

5. Composition As A Checked Normal Form

Arrow induction earns its keep by computing a familiar operation. For $p : x \rightarrow y$, define a motive at (y, p) by the functor category between representables:

$$\text{CompMotive}_x[(y, p)] = \text{Rep}_Z(y) \vdash \text{Rep}_Z(x).$$

At the reflexive outgoing arrow this is $\text{Rep}_Z(x) \vdash \text{Rep}_Z(x)$, so the initial datum is the identity functor. Induction produces a section

$$\text{pathCompSec}(x) = \text{Ind}_x(\text{CompMotive}_x, \text{id}).$$

Its component at p is a functor

$$\text{pathComp}(p) : \text{Rep}_Z(y) \vdash \text{Rep}_Z(x).$$

Applying it to $q : y \rightarrow z$ normalizes to composition:

$$\text{pathComp}(p)[z][q] \rightsquigarrow q \circ p.$$

5.1 The normalization route

The final equation is compact, but the checked route crosses several abstraction boundaries. At a high level,

evaluation performs the following steps:

```

PathInd_transfd(Z)
  ↓ component at x
PathInd_func(Z,x)
  ↓ component at CompMotive_x
path_ind_sec(Z,x,CompMotive_x,id)
  ↓ section component at (y,p)
CompMotive_x[rho_{x,y,p}](id)
  ↓ representable and Sigma-transport
  computation
path_comp_func(p)
  ↓ action at (z,q)
hom_postcomp_fapp0(id_Z,q,p)

```

Each arrow is owned by a reusable projection or computation. The theorem's displayed component projects to fixed-source induction. Section evaluation projects to motive transport along ρ . Sigma transport computes the endpoint from $\text{Rep}_Z(x)[p](\text{id}_x) = p$. The composition motive then exposes postcomposition between representables.

This factorization is valuable for regression. A direct rule `pathComp(p,q) -> q o p` would prove only that one named head reduces. The expanded checks verify that the theorem remains connected to its semantic construction. They instantiate both the displayed telescope route and the derived total-functor route all the way to the same hom-action owner. Nearby negative assertions keep opaque motives, unrelated arrow endpoints, and higher cells from being consumed by the specialized path.

The checks exercise this result both through the primary telescope theorem and through the derived Sigma-total theorem. They also check the preceding representable and ρ computations, so the final result is not an unrelated rewrite attached directly to the theorem name.

6. Usable Categorical Binders

The kernel expressions above are explicit and compositional, but they are not how an end user wants to author every program. Ordinary dependent type theory allows a variable to occur recursively inside a term and reconstructs the wiring. Categorical syntax needs the same usability while distinguishing several kinds of action.

There is a useful implementation subtlety. The runtime normal form is the represented-hom postcomposition owner

```
hom_postcomp_fapp0(id_Z, q, p).
```

The ordinary categorical presentation `comp_fapp0(q,p)` is a typed proof-time comparison surface. Emdash does not globally erase these provenances into one raw composition syntax. Postcomposition, precomposition, and rigid two-endpoint hom action have different variance and different higher behavior. Narrow comparisons relate them when a theorem needs to see the same categorical composite.

This illustrates the normalization policy used throughout the kernel.

- A **rewrite rule** selects a runtime normal form and participates in reduction and critical-pair analysis.
- A **unification rule** helps the framework compare two intended presentations when neither should compute to the other.
- The global `fapp*` and `tapp*` calculus owns generic identity, composition, functoriality, and naturality.
- Constructor-specific rules are projection betas or measured joins, not duplicate copies of generic coherence.

In this sense, the equation $\text{pathComp}(p)(q) = q \circ p$ is both a mathematical theorem and a regression test for the architecture. It passes only if Sigma transport, representable action, displayed transformation projection, section evaluation, and hom-action cut elimination agree on a normal form.

The current surface marks the intrinsic binder mode on the lambda:

```

λ^f    functorial
λ^n    natural / indexed
λ^fd   displayed-functorial
λ^nd   displayed-natural

```

A type annotation such as $x : A$ is separate and may be inferred when the expected type determines it. This is why $\lambda^f x. \dots$, rather than $\lambda x :^f A. \dots$, is the durable notation direction: functoriality belongs to the binder, while the domain annotation is ordinary bidirectional typing information.

Type-directed application. Whitespace application is intentionally neutral in the source. The expression $F x$ does not itself say whether to use an object component, a capped arrow component, a whole hom action, a section evaluation, or a displayed component. The resolver combines the inferred type of F , the classifier and variation of x , and the expected result.

Typed situation	Selected internal reading
$F : A \vdash B, x : \text{Obj}(A)$	object action $\text{fapp0}(F, x)$
$F : A \vdash B, p : \text{Hom}_A(x, y)$	arrow action $\text{fapp1}(F, p)$
functor expected over a whole hom	internal hom action, not one capped point
$s : \Pi(k : ^n K), E[k], k : \text{Obj}(K)$	section component $\text{piapp0}(s, k)$
$FF : E \vdash_K D, a : E[k]$	displayed fibre component
coherent transfor at k	internal component and higher-action owners

The whole-hom case is particularly important. If the expected result is a functor between hom-categories, prematurely choosing the value on one arrow would lose the object needed for the next higher action. The resolver therefore preserves a functor-level term whenever later iteration is requested. This is the surface counterpart of the kernel's preference for functor-level folds over capped point rules.

Contravariance is also typed rather than guessed from names. Represented hom has distinct postcomposition and precomposition owners. A future surface may infer polarity from a classifier or expected mixed-variance family, but the current frontend supports only reviewed cases. It will not reverse an arrow or synthesize an opposite merely because doing so would make a local application typecheck.

Neutral application makes the syntax familiar while keeping the elaboration falsifiable. Each selected route produces an explicit owner that the generic checker can verify. When two routes are possible at the raw syntactic level, the expected type disambiguates them; when the expected type is insufficient, the diagnostic requests an annotation rather than committing to a noncanonical action.

The frontend does not recognize whole source strings and replace them with owner-specific templates. It lowers each scoped callback or parsed binder to a first-order contextual representation and recursively

analyzes subexpressions. Variable occurrence and expected classifier information determine structural wiring.

For ordinary functorial abstraction, the basis includes:

- identity when the body is the variable;
- weakening when the variable is unused;
- diagonal/contraction when it is used more than once;
- exchange for independent nested variables;
- product pairing and projections;
- evaluation of a functor-valued expression against an argument expression;
- functor composition; and
- curry/uncurry for nested abstraction.

This handles examples that are easy to write but not reducible to a single eta pattern:

```

λf x. (H x) (K x)
λf x. F x y0
λf x : A. λf y : B. E y x

```

The first uses x in both function and argument positions. The second abstracts after evaluation at a fixed inner object. The third exchanges independent variables before evaluation. Each compiles to existing product, pairing, evaluation, composition, and exchange owners.

Dependency-aware structural compilation. The compiler treats a context as ordered typed slots plus dependency edges. For each abstraction it first lowers the body recursively, then computes how the selected slot occurs in the resulting typed tree. Zero, one, or multiple uses select weakening, identity-like routing, or contraction/pairing. Independent nested slots may be exchanged. A slot that occurs in the type of a later slot cannot be exchanged across that dependency.

Displayed contexts add a base classifier to every slot. Extending by $a : A[k]$ changes the base from K to $\Sigma_K A$. A later family must name that total or a checked pullback of a family along a known projection. Independent siblings B and C over the same extended base are grouped through $P(B, C)$; after their pair is introduced, the next base is $\Sigma P(B, C)$. The compiler derives these transitions from family types, not from untyped punctuation alone.

This design supports both DTT viewpoints that arise in practice. A sequential telescope represents variables whose types depend on earlier variables. Fibrewise product structure represents several variables that are independent of one another but share a dependent base. Weakening, symmetry, and contraction apply inside the latter block, while substitution and totalization connect the former levels. The implementation may use different lowering routines for the two cases; what matters is that they compose through the same explicit internal owners without an ad hoc semantic exception.

Natural and displayed binders require indexed classifiers rather than an external naturality proof. A section component example is:

```
λ^n k : K. (GG k) ((FF k) (s k))
```

Here $s k$ is a section value in a fibre and $FF k$ is the component of a displayed functor; $GG k$ continues the section through a second displayed functor. Recursive application elaborates this finite rigid chain through generic composition at $Catd_cat K$. Both the object component and the action over a base arrow are already owned by the internal displayed construction.

A displayed functorial example composes internal displayed functors:

```
λ^fd a : E. GG (FF a)
```

Displayed weakening can recover the hidden base index and apply a section:

```
λ^fd a : E. s (index0f a)
```

For independent siblings over one base, the frontend uses fibrewise pairing:

```
λ^fd (b : B, c : C). fibrePair (FF b) (GG c)
```

The mixed telescope combines genuine dependency edges with an independent middle block:

```
λ^fd (a : A; b : B, c : C; d : D).  
  fibrePair b c
```

Semicolons advance the dependent context. The comma says that b and c are siblings over the same preceding total. The family of d is based on the total containing their pair. The contextual compiler derives the necessary pullbacks and totalizations, while the user writes the variables in the same dependency order as an ordinary DTT telescope.

The implemented normal form is not limited to the displayed four-slot example. It accepts any finite sequence of these canonical dependency levels, with finite sibling groups at a level, for the reviewed displayed-functorial and displayed-natural constructions. The category resolver also descends through any finite number of qualified Hom-category levels over its supported roots, and the indexed-section route composes finite rigid chains such as the one above. These results do not license arbitrary dependency or variance graphs, exchange across a dependency edge, or coherence synthesis outside the qualified grammar.

A worked mixed telescope. The mixed example is small enough to run in the reviewer but rich enough to show the dependency algorithm:

```
λ^fd (a : A; b : B, c : C; d : D).  
  fibrePair b c
```

Assume

```
K : Cat
A : Catd K
B, C : Catd (Sigma_cat A)
D : Catd (Sigma_cat P(B,C)).
```

The expected type makes this a displayed contextual functor, not an outer LF lambda or a pointwise function. The compiler extends the base to $\Sigma_K A$, verifies that B and C are independent siblings there, and represents them by the fibrewise product $P(B, C)$. It then checks D over the total containing that pair, inserts dependency-aware weakening because the body does not use d , and emits the sibling projections and displayed pairing owner explicitly. Giving B and C different bases, basing D before the sibling pair, exchanging a across a family that depends on it, or returning a value from an unrelated displayed family is rejected.

The object component alone would not validate this construction. Let p be a base arrow and let u be an internalized arrow in the displayed source. The capped cell of the pairing owner reduces componentwise:

$$\begin{aligned} & \text{cell}(\langle FF, GG \rangle, p, u) \\ &= \langle \text{cell}(FF, p, u), \text{cell}(GG, p, u) \rangle. \end{aligned}$$

The active kernel expresses this at the generic `fdappl_int_cell` projection for `Product_pair_funcd`; it does not introduce a second product-cell calculus. The TypeScript consumer transfers the existing owner and rule, checks the object and internalized-arrow observations, and retains an

opaque cell as a non-collapse witness. Sequential dependency and fibrewise siblings may use different lowering routines; the invariant is that both end in composable internal owners whose object and arrow actions are checked by the same Core.

The displayed-natural slice demonstrates recursive component composition:

```
λ^nd k : K. composeCells (theta k) (eta k)
```

The result is not a pointwise function accompanied by a hand-authored naturality equation. The typed inputs are already transforms, their component operations are internal projections, and generic higher action carries the construction forward.

These examples expose the design's key usability rule:

Convenient syntax may hide structural categorical operations, but it may not fabricate categorical coherence.

If the expected type requests an action for which the active kernel has no internal owner, elaboration reports an unsupported action. Escaped binder tokens, wrong families, incompatible bases, illegal dependency exchange, and foreign construction environments are rejected. This fail-closed policy is what allows the surface to grow without becoming a second, unsound category-theory implementation.

The string syntax is correspondingly bounded. The adapter parses neutral application, annotations, the four binder modes, and the reviewed constructors needed by twelve source presets. It does not parse arbitrary Lambdapi source or every notation in the book. Direct typed TypeScript remains the more complete construction surface; both routes converge before checking.

7. From Surface Terms To Explicit Core

Consider the ordinary source:

```
λ^f x. (H x) (K x)
```

The parser identifies a functorial binder and neutral applications. The resolver uses the expected functor type and the types of `H` and `K` to classify both occurrences of x . The contextual compiler constructs the diagonal, pairs the two resulting functorial expressions, and applies the evaluation functor.

Serialization then reveals the explicit owners and arguments that were silent in the source.

The same path applies to direct TypeScript construction. A scoped builder invokes each callback once with an opaque token. Lowering replaces that token by a locally nameless contextual slot. No arbitrary JavaScript closure enters Core, and a token cannot escape its builder or be reused in another environment.

Explicit Core is then processed by one generic dependent-LF runtime:

1. declarations establish owner types and optional transparent bodies;
2. bidirectional checking verifies applications and dependent binders;

3. metavariable constraints are revisited through bounded pattern unification;

4. β reduces outer lambdas and δ unfolds authorized definitions;

5. reviewed categorical runtime rules normalize internal owners; and

6. the result is serialized with its inferred and expected type.

The strongest small dependent demonstration crosses both layers. An outer LF lambda receives a section over a nested Sigma telescope. Its body applies the section to a dependent pair. Checking verifies that the pair belongs to the expected family; evaluation performs outer β and then the directed Sigma-telescope fibre rule. Replacing the fibre object by one from another family produces a type mismatch rather than a stuck or coerced term.

7.1 Checked behavior, diagnostics, and provenance

Positive examples alone would not establish the boundary. The executable corpus pairs them with negative and non-collapse cases. An object from the wrong category cannot be passed to a functor action. A dependent pair from the wrong family cannot be consumed by a section. A displayed family based on K cannot masquerade as one based on $\Sigma_K A$. An internalized arrow is checked not to collapse to its object component, and an opaque higher cell remains opaque when no projection rule applies.

Diagnostics are normalized at the surface boundary. They record a stable code, phase, source span, expected classifier or family, and an explanatory detail. The browser does not duplicate these checks; it formats the same diagnostic returned by the TypeScript elaborator. This makes an edited example useful to a reviewer: rejection exposes which architectural condition is missing instead of producing a generic JavaScript exception.

Conformance evidence is separate from runtime provenance. A candidate profile records the Lambdapi module, source or canonical-export evidence, selected owner set, and reviewed rule set from which it was built. Node-side tools can emit a deterministic Lambdapi judgment and compare it with the active kernel. The browser consumes only the frozen browser-safe selection contract and compiled runtime data; it neither reads source files nor computes an authority digest.

The browser reviewer packages this architecture without adding semantic machinery. Its expression panel offers twelve presets across $\wedge f$, $\wedge n$, $\wedge fd$, and $\wedge nd$, with editable text, explicit Core, inferred type, structural prerequisites, and source-located rejection. Its evidence panel runs the existing outer-LF, ordinary bracket, and genuinely displayed dependent witnesses on request. Its Core panel exposes a minimal editable dependent-LF script. The entire published runtime is client-side; it opens the overview paper and full book as static assets.

The artifact boundary is worth stating precisely.

Layer	What it establishes
Active Lambdapi sources	Authoritative categorical declarations, computation rules, and proof-time comparisons
Lamdapi checks	Executable positive, negative, and normal-form evidence for the active kernel
TypeScript Core	A small dependent checker/evaluator for the reviewed transferred profile
Categorical frontend	Recursive compilation of the demonstrated binders to internal owners
Browser reviewer	Reproducible access to the same TypeScript paths, with no server or production Lambdapi process

The browser does not contain the entire Lambdapi development. It establishes one real end-to-end path:

surface syntax reaches explicit internal structure, generic checking, and computation in a client an external reviewer can run.

8. From Invertibility Sieves To Local Geometry

Arrow induction began with a family stable under change of endpoint and only then extracted a familiar operation from its total category. The same order organizes local geometry. A condition is first allowed to

vary along every probe. Representability by one chart, coverage by a topology, and reflection into sheaves are separate constructions. Keeping those stages apart is what makes the resulting claims executable rather than mnemonic.

8.1 The sieve comes before the open

Let $\mathcal{O}$ be a presheaf of commutative rings on a category $\mathcal{K}$, let U be an object, and let $s \in \mathcal{O}(U)$. Every probe $p : V \rightarrow U$ restricts the section to $p^*s \in \mathcal{O}(V)$. Define

$$D_U(s)(p) := \text{Unit}_{\mathcal{O}(V)}(p^*s).$$

Ring maps preserve units. Hence a witness at p restricts along every $q : W \rightarrow V$ to a witness at $p \circ q$. The successful probes form an ordinary sieve on U . No decision procedure for invertibility and no representing open were needed. Unit evidence is proposition-valued, so the sieve is subterminal, while a selected inverse remains available during the calculation.

This is the article's organizing geometric reversal: the primary object is invertibility's sieve, not invertibility's open. The familiar phrase “the open on which s is invertible” silently combines two claims. The first is the functorial stability just proved. The second is that all successful probes are represented by one object over U . The first makes sense on an arbitrary site; the second depends on additional geometry.

In a posetal site, a representing open $A \leq U$ satisfies

$$p \in D_U(s) \iff p \text{ factors through } A.$$

It is therefore the largest open on which s is invertible. Zeuner's constructive algebraic geometry [5] organizes the coherent or qcqs case by such compact-open supports. The sieve-centered formulation retains that

account when a compact representative exists, while making the representability question explicit when it does not.

For affine geometry take $\mathbf{Aff} = \mathbf{CRing}^{\text{op}}$. The generalized points of a ring R at a test ring S are maps $R \rightarrow S$. For $f \in R$ the affine invertibility sieve has components

$$D_R(f)(S) = \sum_{h: R \rightarrow S} \mathbf{Unit}_S(h(f)).$$

Now suppose a localization $\iota_f : R \rightarrow R[1/f]$ has been supplied by its universal property. Composition with ι_f sends a map $R[1/f] \rightarrow S$ to a $D_R(f)$ -point. Conversely, a map $h : R \rightarrow S$ carrying f to a unit has a contractible space of factors through ι_f . Selecting its center and using contractibility for uniqueness gives an explicit pointwise equivalence

$$\mathbf{Hom}_{\mathbf{CRing}}(R[1/f], S) \simeq D_R(f)(S).$$

This statement needs no fraction normal form. A quotient, fraction, or fixed-image construction may serve as coordinates once it satisfies the same factorization problem. The checked result is presently pointwise at every test ring; its components retain existing functorial action, but they have not been assembled into a new whole presheaf equivalence.

Multiplication computes intersection at the same semantic level. An image of fg is a unit exactly when the images of both f and g are units, giving $D_R(fg)(S) \simeq D_R(f)(S) \cap D_R(g)(S)$ pointwise. A supplied localization at fg presents the same question without equating independently chosen packages.

8.2 From finite covers to a Cat-valued reflector

The next step is coverage. Suppose a chart ring S carries a finite family $f_1, \dots, f_n$ and coefficients with

$$\sum_i a_i f_i = 1.$$

After selecting localizations at the f_i , the resulting basic chart arrows form a witness-rich proposed cover. The witness retains the family, coefficients, localization packages, and literal containment in the proposed sieve. Coverhood itself should not depend on which presentation was supplied. Emdash therefore defines the generated topology as the intersection of all Grothendieck topologies accepting the generators. It accepts every retained finite presentation, satisfies maximality, pullback stability, and local character, and is least among accepting topologies. The construction proves a universal property; it does not provide an inductive derivation language or a decision procedure for coverhood.

Fix now any site $(\mathcal{K}, J)$ and a Cat-valued presheaf P . For a sieve R on U , whole matching families and whole sections are hom-categories

$$\mathbf{Match}_P(R) = \mathbf{Hom}(\widehat{R}, P),$$

$$\mathbf{Sect}_P(U) = \mathbf{Hom}(yU, P).$$

Restriction is precomposition with the inclusion $\widehat{R} \rightarrow yU$. Locality says that this restriction is an equivalence for every covering sieve. Sheafification asks for more: construct a local presheaf from an arbitrary one, functorially and universally.

The direct cover completion aP is presented by three pieces of categorical structure. **Return** is a whole map $\eta_P : P \rightarrow aP$. **Glue** is a whole functor from matching families in aP to sections in aP , varying functorially over the category of eligible covering questions. **Silent** is one path saying that gluing the restriction of a section recovers that section. Newly glued data may enter later matching families, so glue is recursive; because coefficients are Cat-valued, it must act on arrows between matching families as well as on their objects.

This return/glue/silent pattern is conceptually adapted from Pédrot's free sheaf construction [6], but it is installed directly in categorical semantics. Cover questions, matching and section categories, displayed functoriality, and whole paths live inside the functorial calculus hosted by the outer logical framework. No separate modal type theory is assumed. The

constructors are primitive at this categorical-HIT boundary; their consequences are proved afterward.

Naturality of the one whole glue functor supplies compatibility with pullback of covers. From that compatibility and silent coherence, the opposite restriction-after-glue law is derived rather than postulated. A recursor then extends every seed $P \rightarrow Y$ into a topology-local target Y across aP , and categorical-HIT uniqueness gives the whole equivalence

$$\mathrm{Hom}(aP, Y) \simeq \mathrm{Hom}(P, Y).$$

Consequently direct cover completion assembles a left adjoint

$$a : \mathrm{Psh}_{\mathrm{Cat}}(\mathcal{K}) \rightleftarrows \mathrm{Sh}_{\mathrm{Cat}}(\mathcal{K}, J) : i, \quad a \dashv i,$$

whose counit on every sheaf is an omega-equivalence. This is a fixed-site, Cat-valued reflector. It does not yet lift the construction to commutative-ring-valued presheaves, prove left exactness, or supply a base-change theorem.

8.3 The precise boundary of the geometric application

The affine chain now has a constructed computational spine. Ring maps give generalized probes; unit preservation gives $D_R(f)$; localization represents that question pointwise; products compute intersections; finite unimodular families generate a big Zariski topology; and the coordinate presheaf computes restriction along its charts. The Cat-valued reflector demonstrates that sheafification itself can be expressed directly at the layer of categorical semantics.

The remaining commitments are deliberately visible. The active affine presentation supplies a reflective commutative-ring-valued structure sheaf and whole localization locality rather than deriving them from the Cat-valued reflector. A global-first, site-relative scheme presentation begins with one supplied global ringed

object, a covering sieve constructively generated by two selected affine charts, and topology-local ring behavior. Whole restrictions and a selected chart intersection are inherited from the global presheaf instead of copied into an atlas record.

On that literal overlap, polynomial and localization universal properties construct the Laurent coordinate changes $t \mapsto u^{-1}$ and $u \mapsto t^{-1}$. A supplied projective-line package retains the global site-relative scheme, its actual selected overlap, and the whole Laurent comparison. It is an end-to-end computational capability, not a construction of a global object from abstract charts. There is no graded ring interface, homogeneous localization, degree-zero construction, `Proj`, general projective space, or non-affineness theorem in the active artifact.

8.4 Free groupoidal realization and the Gray test

The same owner discipline connects directed arrows with equality paths. For a groupoidal classifier A , the category $\mathrm{Path}(A)$ has elements of A as objects and equality as hom. Products are closed under this view homwise: a path of pairs splits into its two coordinates and coordinate paths reassemble, while direct dependent transport agrees propositionally with either sequential coordinate order.

The opaque Circle makes the computational boundary sharper. Its dependent eliminator reduces both at the base point and when dependent path action is observed on the generating loop. The familiar constant-family `ap` equation is derived propositionally rather than installed

as a second runtime rule. A successor-localized Integer classifier supplies the universal cover, and encode/decode proves

$$\mathrm{Hom}_{S^1}(\mathrm{base}, \mathrm{base}) \simeq \mathbb{Z}.$$

WalkingEnd maps to the Circle by sending its directed generator to the loop; natural powers become the nonnegative integer powers. The two-ended WalkingArrow similarly maps to a groupoidal Interval. These are finite tests of the category-indexed operation $\mathrm{Groupoidify}(C)$, whose constructor is one whole functor

$$\eta_C : C \longrightarrow \mathrm{Path}(\mathrm{Groupoidify}(C)).$$

For every groupoid G , restriction along η_C and the whole extension recursor form a fixed-forward mapping-object equivalence

$$\mathrm{Hom}_{\mathrm{Grpd}}(\mathrm{Groupoidify}(C), G)$$

$$\simeq_{\omega} \mathrm{Functor}(C, \mathrm{Path}(G)).$$

The recursor computes on represented points and dependent first cells. Restriction and extension have whole beta/eta paths and retain higher action. This is more than a carrier-level free groupoid, but less than a packaged adjunction: the action of `Groupoidify` on a source functor has not yet been constructed.

The directed side of the same story is exposed by a compositor

$$\phi_{g,f}^F : F[g] \circ F[f] \Longrightarrow F[g \circ f]$$

projected from whole internal laxity. In a path target this cell is invertible; strictness identifies it with an equality-induced arrow; in an arbitrary directed target it

may remain noninvertible. Emdash reuses that distinction in the profiled category $\mathrm{GrayHom}_{\mathrm{lax}}(A, B)$: objects are functors paired with their strictness property, while arrows and all higher Homs come from the ambient transfor tower.

One selected right closure is checked:

$$\mathrm{GrayHom}_{\mathrm{lax}}(A \otimes_R B, C)$$

$$\simeq_{\omega} \mathrm{GrayHom}_{\mathrm{lax}}(A, \mathrm{GrayHom}_{\mathrm{lax}}(B, C)).$$

Coevaluation at two walking arrows gives a four-vertex square. Its two coordinate routes are compared by a nonidentity interchanger projected from whole laxity, with one next action retained. This is a computational coherence stress test, not a claim to the full combinatorial or Crans--Gray tensor studied in [7]. Together, the geometry and groupoidal slices support the same architectural conclusion: generic action should own functoriality and naturality, semantic constructors should own computation, and elaboration should reconstruct structural wiring.

8.5 Simplexes from dependent homs

Injective skip/keep codes form the internal augmented semi-simplex category. Yoneda gives $\Delta[n] = \mathrm{Hom}(-, n + 1)$. Directed join gives $\Delta[0] = \mathbf{1}$ and $\Delta[n + 1] = \Delta[n] \star \mathbf{1}$; the native recurrence is $S_0(C) = C$ and $S_{k+1} = \mathrm{PathOut}_{S_k}(s_k)$.

Since `PathOut` is a Sigma of a representable hom, its arrows pair a base cell with a dependent cell above transport. A triangle contains $p_{12} \circ p_{01} \Longrightarrow p_{02}$; whole base and endpoint action supply the extra tetrahedral faces and retain another hom action.

An intrinsic flag code records the changing native category. A stage $F, G : K \rightarrow B$, $\epsilon : F \Rightarrow G$

sends an old source s to the new code `step(code, F[s])` and source $(G[s], \epsilon[s])$. The first stage comes from the ordinal join and later stages lift ϵ through `PathOut`. Nat recursion therefore constructs the source for variable n , maps it under every $H : \Delta[n] \rightarrow C$, and restricts it along nonempty faces. Dimensions zero through four, the five faces of the four-simplex, noncollapse, and one next action are checked.

The active `DependentSimplexObservation(C, n)` is an object package, not yet a whole category. The corresponding mapping-category equivalence, degeneracies, and general Kan, Segal, or Rezk theorems remain open.

9. Research Boundaries

The word *checked* in this article has a local meaning. `Lambdapi` accepts the active declarations and rules and

verifies the recorded assertions; the TypeScript runtime accepts the demonstrated transferred terms and rejects

its negative corpus. It does not mean that a global metatheory has already been formalized.

The principal open boundaries are:

- **Surface generality.** The categorical text language covers the reviewed twelve-preset profile, not arbitrary book notation, Lambdapi syntax, or every direct TypeScript constructor. The full canonical mathematical surface is intentionally broader.
- **Displayed depth and variance.** Independent siblings, genuine dependency, finite canonical dependency levels and sibling groups, qualified finite Hom-category recursion, finite rigid section chains, and selected higher action are implemented. Arbitrary dependency or variance graphs, dependency-edge exchange, unrestricted mixed introduction and currying, polarity-directed contravariant lowering, and general displayed-transfor coherence are not.
- **Transfer scale.** Generic transfer engines and varied representative tranches exist, but systematic transfer of the whole active Lambdapi library has not graduated. Mathematical rules will continue to require owner and interaction review even if their representation becomes mechanical.
- **Sheaffication and geometry.** Direct cover completion constructs a fixed-site Cat-valued reflector, not a commutative-ring lift, a left-exactness theorem, or base-change semantics. Affine structure-sheaf and locality capabilities remain supplied; schemes are site-relative; the projective-line package is supplied rather than glued from charts; and graded `Proj` and projective space remain future work.
- **Categorical closure.** General dependent adjunctions $\Sigma_F \dashv F^* \dashv \Pi_F$, full profunctor tensor/coend semantics, complete equipment coherence, and semantic collage or dependent elimination for join remain future work.
- **Groupoidal closure.** Equality, J, groupoidal dependent sums/products, the Circle HIT with $\Omega S^1 \simeq \mathbb{Z}$, its concrete `WalkingEnd` comparison, the `WalkingArrow/Interval` test, category-indexed `Groupoidify(C)` with its target-side whole mapping equivalence, and representative product preservation/transport coherence are checked. Source functoriality and the adjunction package, closure for every former, and arbitrary motive-directed computational decomposition of J are not.
- **Metatheory.** No global normalization, confluence, canonicity, consistency, decidability, or semantic-soundness theorem is claimed for the full combined calculus. Lambdapi's local rule checks and the project's diagnostics are implementation evidence, not replacements for those theorems.
- **Higher categories.** Whole laxity, strict-functor property packages, one profiled right Gray closure, and its walking interchanger are checked. The mirror closure, tensor functoriality and coherence, full `Crans--Gray` monoidality, global migration of historical strict endpoint cuts, and a general weak- ω -category metatheory remain open.
- **Simplicial methods.** Injective face codes, the augmented semi-simplex index, native simplexes through dimension four, and the variable canonical source are checked. Degeneracies, the whole mapping-category equivalence, and general Kan, Segal, Rezk, complicial, or oriental structure are not.

These limits are architectural information, not disclaimers added after the fact. Fail-closed elaboration, explicit product manifests, runtime versus proof-time ownership, and the separation of authority from conformance exist to make partial progress reviewable without silently broadening the claim.

10. Conclusion

Functorial Type Theory treats categorical variation as computational structure. Emdash realizes that idea through complementary layers: an outer dependent logical framework and an inner directed dependent calculus of families, totals, sections, dependent homs, and higher action.

Synthetic arrow induction exhibits the mathematical payoff. The canonical Sigma-transport arrow ρ turns reflexive data into a section over all outgoing arrows, and the composition motive normalizes to $q \circ p$. Usable categorical binders exhibit the implementation payoff. Recursive ordinary and displayed variable occurrences compile to explicit internal owners and are checked by one generic TypeScript LF, including object and arrow behavior.

Sieve-centered geometry exhibits the local-to-global payoff. Invertibility is first the stable sieve $D_R(f)$, a localization represents that question pointwise, finite certified charts generate a topology, and direct cover completion constructs a Cat-valued reflector through whole return, glue, silent coherence, recursion, and uniqueness. The boundaries beyond that reflector remain as explicit as the construction itself.

Groupoidal realization exhibits the higher-dimensional payoff. Circle encode/decode restores the inverse powers absent from WalkingEnd, generic groupoidification characterizes maps out by a whole equivalence, and the Gray walking square retains a

directed interchanger rather than erasing it into an equation.

Dependent simplexes exhibit the recursive payoff: face observations and higher action arise from iterated outgoing paths rather than an expanding list of coherence fields.

The current system is bounded, but the design question has a concrete answer: readable binders, explicit Core, checked categorical computation, an authoritative Lambdapi kernel, sieve-based local geometry, free groupoidal realization, and a client-side reviewer fit into one architecture. The remaining work is to extend its mathematical and transfer coverage without losing that internalized, normalization-first discipline.

References

1. The Univalent Foundations Program. *Homotopy Type Theory: Univalent Foundations of Mathematics*. Institute for Advanced Study, 2013.
2. Kosta Došen. *Cut Elimination in Categories*. Trends in Logic 6. Kluwer Academic Publishers, 1999.
3. The Lambdapi contributors. *Lamdapi User Manual*. lambdapi.readthedocs.io.
4. The emdash contributors. *emdash v3.2 sources*. Accompanying artifact. [Code repository](#).
5. Max Zeuner. *Univalent Foundations of Constructive Algebraic Geometry*. arXiv:2407.17362v1, 2024.
6. Pierre-Marie Pédro. “Pursuing Shtuck.” Preprint, 2023. HAL: hal-04251754v1.
7. Amar Hadzahasanovic. *Combinatorics of Higher-Categorical Diagrams*. arXiv:2404.07273v2, 2024.
8. The emdash contributors. *Functorial Type Theory: Univalent Foundations for Mathematics*. Ongoing book, 2026. [Book DOI 10.5281/zenodo.21544186](#). [Code and source repository](#).